

Image Processing

SOS CS11

Mid-Term Report

Panav Shah
23B3323

June 2024

Contents

1	Diving into Deep Learning	3
1.1	Deep Learning in Image Processing	3
1.1.1	Image Classification	3
1.1.2	Object Detection	3
1.1.3	Image Segmentation	3
1.1.4	Image Enhancement	3
1.2	Key concepts in Deep Learning for Image Processing	3
1.3	Tools and Frameworks	4
2	Image Classification	4
2.1	Classification of FashionMNIST Dataset	4
2.1.1	Importing computer vision libraries	4
2.1.2	Getting the dataset	5
2.1.3	Visualise our data	5
2.1.4	Prepare DataLoader	6
2.1.5	Building a Convolutional Neural Network (CNN)	6
2.1.6	Setting up Loss function, Optimizer and Accuracy function	8
2.1.7	Defining each train and test step	8
2.1.8	Training and Testing our model using the <i>train_{step}</i> and <i>test_{step}</i> functions	10
2.1.9	Evaluate random predictions using our model	11
2.1.10	Plot a Confusion Matrix	13
2.2	Classification of a Pizza, Steak and Sushi (a subset of the Food101 dataset)	14
2.2.1	Importing computer vision libraries, setting up device agnostic code and importing the dataset	14
2.2.2	Transforming data	15
2.2.3	Loading images to a dataset	17
2.2.4	Turn loaded images into DataLoader's	18
2.2.5	Create a TinyVGG model	19
2.2.6	Create train step, test step and train functions	20
2.2.7	Train and evaluate model ₀	23
2.2.8	Making a prediction on a custom image	23
3	Bibliography	25

1 Diving into Deep Learning

Deep learning is a method in artificial intelligence (AI) that teaches computers to process data in a way that is inspired by the human brain.

1.1 Deep Learning in Image Processing

Deep learning has revolutionized the field of image processing by providing powerful techniques for tasks such as image classification, object detection, segmentation, and image generation. Here's an overview of how deep learning is applied in image processing:

1.1.1 Image Classification

Convolutional Neural Networks (CNNs): CNNs are the backbone of image classification tasks. They use layers of convolutional filters to extract features from images, followed by fully connected layers that classify the images based on the extracted features. Famous architectures include AlexNet, VGGNet, ResNet, and Inception.

1.1.2 Object Detection

Region-Based CNNs (R-CNN): The original goal of R-CNN was to take an input image and produce a set of bounding boxes as output, where each bounding box contains an object and also the category (e.g. car or pedestrian) of the object. More recently, R-CNN has been extended to perform other computer vision tasks. The following covers some of the versions of R-CNN that have been developed.

1.1.3 Image Segmentation

Semantic Segmentation: Classify the object class for each pixel within an image. That means there is a label for each pixel.

1.1.4 Image Enhancement

Super-Resolution: Increasing the resolution of images using models like Super-Resolution CNN (SRCNN) and Generative Adversarial Networks for Image Super-Resolution (SRGAN).

1.2 Key concepts in Deep Learning for Image Processing

- **Convolution :** Convolution is a general purpose filter effect for images. It is a matrix applied to an image and a mathematical operation, comprised of integers

- **Pooling :** The pooling operation involves sliding a two-dimensional filter over each channel of feature map and summarising the features lying within the region covered by the filter.
- **Activation Functions :** Non-linear functions like ReLU (Rectified Linear Unit) that introduce non-linearity into the model, allowing it to learn more complex patterns.
- **Backpropogation :** The process of updating the model's weights by calculating the gradient of the loss function with respect to each weight and adjusting the weights to minimize the loss.
- **Data Augmentation :** Techniques such as rotation, scaling, flipping, colour variation, ect. used to artificially increase the size of the training dataset, improving the model's accuracy.

1.3 Tools and Frameworks

PyTorch : PyTorch is a machine learning library based on the Torch library, used for applications such as computer vision and natural language processing, originally developed by Meta AI and now part of the Linux Foundation umbrella. It is recognized as one of the two most popular machine learning libraries alongside TensorFlow, offering free and open-source software released under the modified BSD license. Although the Python interface is more polished and the primary focus of development, PyTorch also has a C++ interface.

2 Image Classification

2.1 Classification of FashionMNIST Dataset

2.1.1 Importing computer vision libraries

```

1 # Import PyTorch
2 import torch
3 from torch import nn
4
5 # Import torchvision
6 import torchvision
7 from torchvision import datasets
8 from torchvision import transforms
9 from torchvision.transforms import ToTensor
10
11 # Import matplotlib for visualisation
12 import matplotlib.pyplot as plt
13
14 # Check versions
15 print(torch.__version__)
16 print(torchvision.__version__)

```

2.3.0+cu121
0.18.0+cu121

2.1.2 Getting the dataset


```
1 # Setup training data
2 from torchvision import datasets
3 train_data = datasets.FashionMNIST(
4     root="data", # where to download the data to?
5     train=True, # do we want the training dataset?
6     download=True, # do we want to download yes/no?
7     transform=torchvision.transforms.ToTensor(), # how do we want
8     to transform the data?
9     target_transform=None # how do we want to transform the labels/
10     targets?
11 )
12
13 test_data = datasets.FashionMNIST(
14     root="data", # where to download the data to?
15     train=False, # do we want the training dataset?
16     download=True, # do we want to download yes/no?
17     transform=torchvision.transforms.ToTensor(), # how do we want
18     to transform the data?
19     target_transform=None # how do we want to transform the labels/
20     targets?
21 )
22
23 image, label = train_data[0]
24 class_names = train_data.classes
25 class_to_idx = train_data.class_to_idx
26
27 print(f"Image shape: {image.shape} -> [color_channels, height,
28     width]")
29 print(f"Image label: {class_names[label]}")
```

Image shape: torch.Size([1, 28, 28]) -> [color_channels, height, width]
Image label: Ankle boot

2.1.3 Visualise our data

```
1 plt.imshow(image.squeeze(), cmap="gray")
2 plt.title(class_names[label])
```

Text(0.5, 1.0, 'Ankle boot')



2.1.4 Prepare DataLoader

```

1 from torch.utils.data import DataLoader
2
3 # Setup the batch size hyperparameter
4 BATCH_SIZE = 32
5
6 # Turn datasets into iterables (batches)
7 train_dataloader = DataLoader(dataset=train_data,
8                               batch_size=BATCH_SIZE,
9                               shuffle=True)
10
11 test_dataloader = DataLoader(dataset=test_data,
12                              batch_size=BATCH_SIZE,
13                              shuffle=False)
14
15 print(f"Length of train_dataloader: {len(train_dataloader)} batches
16       of {BATCH_SIZE}")
17 print(f"Length of test_dataloader: {len(test_dataloader)} batches
18       of {BATCH_SIZE}")

```

Length of train_dataloader: 1875 batches of 32
Length of test_dataloader: 313 batches of 32

2.1.5 Building a Convolutional Neural Network (CNN)

```

1 # Create a convolutional neural network
2 class FashionMNISTModel(nn.Module):
3     """
4     Model architecture that replicates the TinyVGG
5     model from CNN explainer website.
6     """

```

```

7  def __init__(self, input_shape: int, hidden_units: int,
8      output_shape: int):
9      super().__init__()
10     self.conv_block_1 = nn.Sequential(
11         nn.Conv2d(in_channels=input_shape,
12                   out_channels=hidden_units,
13                   kernel_size=3,
14                   stride=1,
15                   padding=1), # values we can set ourselves in our
16         NN's are called hyperparameters
17         nn.ReLU(),
18         nn.Conv2d(in_channels=hidden_units,
19                   out_channels=hidden_units,
20                   kernel_size=3,
21                   stride=1,
22                   padding=1),
23         nn.ReLU(),
24         nn.MaxPool2d(kernel_size=2)
25     )
26     self.conv_block_2 = nn.Sequential(
27         nn.Conv2d(in_channels=hidden_units,
28                   out_channels=hidden_units,
29                   kernel_size=3,
30                   stride=1,
31                   padding=1), # values we can set ourselves in our
32         NN's are called hyperparameters
33         nn.ReLU(),
34         nn.Conv2d(in_channels=hidden_units,
35                   out_channels=hidden_units,
36                   kernel_size=3,
37                   stride=1,
38                   padding=1),
39         nn.ReLU(),
40         nn.MaxPool2d(kernel_size=2)
41     )
42     self.classifier = nn.Sequential(
43         nn.Flatten(),
44         nn.Linear(in_features=hidden_units*7*7, # There is a trick
45                   to calculate this
46                   out_features=output_shape)
47     )
48
49     def forward(self, x):
50         x = self.conv_block_1(x)
51         x = self.conv_block_2(x)
52         x = self.classifier(x)
53         return x
54
55 model = FashionMNISTModel(input_shape=1,
56                           hidden_units=10,
57                           output_shape=len(class_names)).to(
58     device)
59 model

```

```

FashionMNISTModelV2(
  (conv_block_1): Sequential(
    (0): Conv2d(1, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU()
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (conv_block_2): Sequential(
    (0): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU()
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (classifier): Sequential(
    (0): Flatten(start_dim=1, end_dim=-1)
    (1): Linear(in_features=490, out_features=10, bias=True)
  )
)

```

2.1.6 Setting up Loss function, Optimizer and Accuracy function

```

1 loss_fn = nn.CrossEntropyLoss()
2 optimizer=torch.optim.SGD(params=model_2.parameters(),
3                             lr=0.1)
4 def accuracy_fn(y_true, y_pred):
5     correct = torch.eq(y_true, y_pred).sum().item()
6     acc = (correct / len(y_pred)) * 100
7     return acc

```

2.1.7 Defining each train and test step

```

1 def train_step(model: torch.nn.Module,
2                 data_loader: torch.utils.data.DataLoader,
3                 loss_fn: torch.nn.Module,
4                 optimizer: torch.optim.Optimizer,
5                 accuracy_fn,
6                 device: torch.device = device):
7     """ Performs a training with model trying to learn on data_loader
8     """
9     train_loss, train_acc = 0, 0
10    model.to(device)
11
12    # Put model into training mode
13    model.train()
14
15    # Add a loop to loop through the training batches
16    for batch, (X, y) in enumerate(data_loader):
17        # Put data on target device

```



```

17 X, y = X.to(device), y.to(device)
18
19 # 1. Forward pass (outputs the raw logits from the model)
20 y_pred = model(X)
21
22 # 2. Calculate loss and accuracy (per batch)
23 loss = loss_fn(y_pred, y)
24 train_loss += loss # accumulate train loss
25 train_acc += accuracy_fn(y_true=y,
26                           y_pred=y_pred.argmax(dim=1)) # go from
    logits - prediction labels
27
28 # 3. Optimizer zero grad
29 optimizer.zero_grad()
30
31 # 4. Loss backward
32 loss.backward()
33
34 # 5. Optimizer step
35 optimizer.step()
36
37
38 # Divide total train loss by length of train dataloader
39 train_loss /= len(data_loader)
40 train_acc /= len(data_loader)
41 print(f"Train loss: {train_loss:.5f} | Train acc: {train_acc:.2f
    }%")

```

```

1 def test_step(model: torch.nn.Module,
2               data_loader: torch.utils.data.DataLoader,
3               loss_fn: torch.nn.Module,
4               accuracy_fn,
5               device: torch.device = device):
6     """Performs a testing loop step on model going over data_loader.
7     """
8     test_loss, test_acc = 0, 0
9     model.to(device)
10    model.eval()
11
12    # Turn on inference mode context manager
13    with torch.inference_mode():
14        for X, y in data_loader:
15            # Send the data to the target device
16            X, y = X.to(device), y.to(device)
17
18            # 1. Forward pass
19            test_pred = model(X)
20
21            # 2. Calculate the loss/acc
22            test_loss += loss_fn(test_pred, y)
23            test_acc += accuracy_fn(y_true=y,
24                                    y_pred=test_pred.argmax(dim=1)) # go
from logits -> prediction labels
25        # Adjust metrics and print out
26        test_loss /= len(data_loader)
27        test_acc /= len(data_loader)

```

```

28     print(f"Test loss: {test_loss:.5f} | Test acc: {test_acc:.2f}%\n")

```

2.1.8 Training and Testing our model using the `train_step` and `test_step` functions

```

1 # Train and test model
2 epochs = 10
3 for epoch in range(epochs):
4     print(f"Epoch: {epoch}\n-----")
5     train_step(model=model,
6                 data_loader=train_dataloader,
7                 loss_fn=loss_fn,
8                 optimizer=optimizer,
9                 accuracy_fn=accuracy_fn,
10                device=device)
11    test_step(model=model,
12              data_loader=test_dataloader,
13              loss_fn=loss_fn,
14              accuracy_fn=accuracy_fn,
15              device=device)

```

```

Epoch: 0
-----
Train loss: 0.59531 | Train acc: 78.43%
Test loss: 0.39706 | Test acc: 85.53%

```

```

Epoch: 1
-----
Train loss: 0.36254 | Train acc: 87.01%
Test loss: 0.34667 | Test acc: 87.16%

```

```

Epoch: 2
-----
Train loss: 0.32573 | Train acc: 88.22%
Test loss: 0.31780 | Test acc: 88.42%

```

```

Epoch: 3
-----
Train loss: 0.30446 | Train acc: 88.88%
Test loss: 0.32878 | Test acc: 87.88%

```

```

Epoch: 4
-----
Train loss: 0.29002 | Train acc: 89.44%
Test loss: 0.30029 | Test acc: 89.22%

```

```

Epoch: 5
-----

```

Train loss: 0.27894 | Train acc: 89.83%
Test loss: 0.30936 | Test acc: 89.16%

Epoch: 6

Train loss: 0.26996 | Train acc: 90.23%
Test loss: 0.31740 | Test acc: 88.49%

Epoch: 7

Train loss: 0.26435 | Train acc: 90.36%
Test loss: 0.29987 | Test acc: 89.01%

Epoch: 8

Train loss: 0.25912 | Train acc: 90.54%
Test loss: 0.30351 | Test acc: 89.90%

Epoch: 9

Train loss: 0.25476 | Train acc: 90.76%
Test loss: 0.30208 | Test acc: 89.19%

2.1.9 Evaluate random predictions using our model

```
1 def make_predictions(model: torch.nn.Module,
2                       data: list,
3                       device: torch.device = device):
4     pred_probs = []
5     model.eval()
6     with torch.inference_mode():
7         for sample in data:
8             # Prepare the sample (add a batch dimension and pass to
9             # target device)
10            sample = torch.unsqueeze(sample, dim=0).to(device)
11
12            # Forward pass (model outputs raw logits)
13            pred_logit = model(sample)
14
15            # Get prediction probability (logit -> prediction probability)
16            pred_prob = torch.softmax(pred_logit.squeeze(), dim=0)
17
18            # Get pred_prob off the GPU for further calculations
19            pred_probs.append(pred_prob.cpu())
20
21        # Stack the pred_probs to turn list into a tensor
22        return torch.stack(pred_probs)
23
24 # Make predictions
25 pred_probs = make_predictions(model=model,
```

```

25         data=test_samples)
26
27 # Convert prediction probabilities to labels
28 pred_classes = pred_probs.argmax(dim=1)
29 pred_classes
30
31 # Plot predictions
32 plt.figure(figsize=(9, 9))
33 nrows = 3
34 ncols = 3
35 for i, sample in enumerate(test_samples):
36     # Create subplot
37     plt.subplot(nrows, ncols, i+1)
38
39     # Plot the target image
40     plt.imshow(sample.squeeze(), cmap="gray")
41
42     # Find the prediction (in text form, e.g. "Sandal")
43     pred_label = class_names[pred_classes[i]]
44
45     # Get the truth label (in text form)
46     truth_label = class_names[test_labels[i]]
47
48     # Create a title for the plot
49     title_text = f"Pred: {pred_label} | Truth: {truth_label}"
50
51     # Check for equality between pred and truth and change color of
52     # the title text
53     if pred_label == truth_label:
54         plt.title(title_text, fontsize=10, c="g")
55     else:
56         plt.title(title_text, fontsize=10, c="r")
57
58     plt.axis(False)

```



2.1.10 Plot a Confusion Matrix

```


1 # 1. Make predictions with trained model
2 y_preds = []
3 model_2.eval()
4 with torch.inference_mode():
5     for X, y in test_dataloader:
6         # Send the data and targets to target device
7         X, y = X.to(device), y.to(device)
8         # Do the forward pass
9         y_logit = model_2(X)
10        # Turn predictions from logits -> prediction probabilities ->
        prediction labels
11        y_pred = torch.softmax(y_logit.squeeze(), dim=0).argmax(dim=1)
12        # Put predictions on CPU for evaluation
13        y_preds.append(y_pred.cpu())
14
15 # Concatenate list of predictions into a tensor
16 y_pred_tensor = torch.cat(y_preds)
17
18 from torchmetrics import ConfusionMatrix
19 from mlxtend.plotting import plot_confusion_matrix
20

```

```

21 # 2. Setup confusion instance and compare predictions to targets
22 confmat = ConfusionMatrix(num_classes=len(class_names), task="
    multiclass")
23 confmat_tensor = confmat(preds=y_pred_tensor,
24                           target=test_data.targets)
25
26 # 3. Plot the confusion matrix
27 fig, ax = plot_confusion_matrix(
28     conf_mat=confmat_tensor.numpy(), # matplotlib likes working
    with numpy
29     class_names=class_names,
30     figsize=(10, 7)
31 )

```



2.2 Classification of a Pizza, Steak and Sushi (a subset of the Food101 dataset)

2.2.1 Importing computer vision libraries, setting up device agnostic code and importing the dataset

```

1 import torch

```

```

2 from torch import nn
3
4 # Setup device-agnostic code
5 device = "cuda" if torch.cuda.is_available() else "cpu"
6
7 import requests
8 import zipfile
9 from pathlib import Path
10
11 # Setup path to a data folder
12 data_path = Path("data/")
13 image_path = data_path / "pizza_steak_sushi"
14
15 # If the image folder doesn't exist, download it and prepare it...
16 if image_path.is_dir():
17     print(f"{image_path} directory already exists... skipping
18         download")
19 else:
20     print(f"{image_path} does not exist... creating one")
21     image_path.mkdir(parents=True, exist_ok=True)
22
23 # Download pizza, steak and sushi data
24 with open(data_path / "pizza_steak_sushi.zip", "wb") as f:
25     request = requests.get("https://github.com/mrdbourke/pytorch-
26         deep-learning/blob/main/data/pizza_steak_sushi.zip?raw=true")
27     print("Downloading pizza, steak, sushi data...")
28     f.write(request.content)
29
30 # Unzip pizza, steak, sushi data
31 with zipfile.ZipFile(data_path / "pizza_steak_sushi.zip", "r") as
32     zip_ref:
33     print("Unzipping pizza, steak and sushi data...")
34     zip_ref.extractall(image_path)

```

data/pizza_steak_sushi does not exist... creating one
 Downloading pizza, steak, sushi data...
 Unzipping pizza, steak and sushi data...

2.2.2 Transforming data

```


1 import torch
2 from torch.utils.data import DataLoader
3 from torchvision import datasets, transforms
4
5 # Write a transform for image
6 data_transform = transforms.Compose([
7     # Resize our images to 64x64
8     transforms.Resize(size=(64, 64)),
9     # Flip the images randomly on the horizontal
10    transforms.RandomHorizontalFlip(p=0.5),
11    # Turn the image into a torch.tensor
12    transforms.ToTensor()
13 ])
14

```

```

15 def plot_transformed_images(image_path: list, transform, n=3, seed
    =42):
16     """
17     Selects random images from a path of images and loads/transforms
        them
18     then plots the original vs the transformed version.
19     """
20     if seed:
21         random.seed(seed)
22     random_image_paths = random.sample(image_path_list, k=n)
23     for image_path in random_image_paths:
24         with Image.open(image_path) as f:
25             fig, ax = plt.subplots(nrows=1, ncols=2)
26             ax[0].imshow(f)
27             ax[0].set_title(f"Original\nSize: {f.size}")
28             ax[0].axis(False)
29
30             # Transform and plt target image
31             transformed_image = transform(f).permute(1, 2, 0) # note we
        will need to change shape
32             ax[1].imshow(transformed_image)
33             ax[1].set_title(f"Transformed\nShape: {transformed_image.
        shape}")
34             ax[1].axis("off")
35
36             fig.suptitle(f"Class: {image_path.parent.stem}", fontsize=16)
37
38 plot_transformed_images(image_path=image_path_list,
39                         transform=data_transform,
40                         n=2,
41                         seed=42)

```

2.2.3 Loading images to a dataset

```

1 # Use ImageFolder to create datasets
2 from torchvision import datasets
3 train_data = datasets.ImageFolder(root=train_dir,
4                                   transform=data_transform, # a
5                                   target_transform=None) # a
6     transform for the data
7     transform for the label
8
9 test_data = datasets.ImageFolder(root=test_dir,
10                                  transform=data_transform,
11                                  target_transform=None)
12
13 class_names = train_data.classes
14 class_dict = train_data.class_to_idx

```

```

14
15 img, label = train_data[0]
16 print(f"Image tensor:\n {img}")
17 print(f"Image shape: {img.shape}")
18 print(f"Image datatype: {img.dtype}")
19 print(f"Image label: {label}")
20 print(f"Label datatype: {type(label)}")

Image tensor:
tensor([[[[0.1137, 0.1020, 0.0980, ..., 0.1255, 0.1216, 0.1176],
          [0.1059, 0.0980, 0.0980, ..., 0.1294, 0.1294, 0.1294],
          [0.1020, 0.0980, 0.0941, ..., 0.1333, 0.1333, 0.1333],
          ...,
          [0.1098, 0.1098, 0.1255, ..., 0.1686, 0.1647, 0.1686],
          [0.0902, 0.0941, 0.1098, ..., 0.1686, 0.1647, 0.1686],
          [0.0863, 0.0863, 0.0980, ..., 0.1686, 0.1647, 0.1647]],
        [[0.0745, 0.0706, 0.0745, ..., 0.0588, 0.0588, 0.0588],
          [0.0745, 0.0706, 0.0745, ..., 0.0627, 0.0627, 0.0627],
          [0.0706, 0.0745, 0.0745, ..., 0.0706, 0.0706, 0.0706],
          ...,
          [0.1255, 0.1333, 0.1373, ..., 0.2510, 0.2392, 0.2392],
          [0.1098, 0.1176, 0.1255, ..., 0.2510, 0.2392, 0.2314],
          [0.1020, 0.1059, 0.1137, ..., 0.2431, 0.2353, 0.2275]],
        [[0.0941, 0.0902, 0.0902, ..., 0.0157, 0.0196, 0.0196],
          [0.0902, 0.0863, 0.0902, ..., 0.0196, 0.0157, 0.0196],
          [0.0902, 0.0902, 0.0902, ..., 0.0157, 0.0157, 0.0196],
          ...,
          [0.1294, 0.1333, 0.1490, ..., 0.1961, 0.1882, 0.1843],
          [0.1098, 0.1137, 0.1255, ..., 0.1922, 0.1843, 0.1804],
          [0.1059, 0.0980, 0.1059, ..., 0.1882, 0.1804, 0.1765]]]])

Image shape: torch.Size([3, 64, 64])
Image datatype: torch.float32
Image label: 0
Label datatype: <class 'int'>

```

2.2.4 Turn loaded images into DataLoader's

```

1 # Turn train and test datasets into DataLoader's
2 from torch.utils.data import DataLoader
3 BATCH_SIZE = 1
4 train_dataloader = DataLoader(dataset=train_data,
5                               batch_size=BATCH_SIZE,
6                               shuffle=True)
7 test_dataloader = DataLoader(dataset=test_data,
8                              batch_size=BATCH_SIZE,
9                              shuffle=False)

```

2.2.5 Create a TinyVGG model

```
1 # Create simple transform
2 simple_transform = transforms.Compose([
3     transforms.Resize(size=(64, 64)),
4     transforms.ToTensor()
5 ])
6
7 # 1. Load and transfer data
8 from torchvision import datasets
9 train_data_sample = datasets.ImageFolder(root=train_dir,
10                                         transform=simple_transform
11 )
12 test_data_sample = datasets.ImageFolder(root=test_dir,
13                                         transform=simple_transform
14 )
15
16 # 2. Turn the datasets into Dataloaders
17 import os
18 from torch.utils.data import DataLoader
19
20 # Setup batch size and number of workers
21 BATCH_SIZE = 32
22 NUM_WORKERS = os.cpu_count()
23
24 class TinyVGG(nn.Module):
25     """
26     Model architecture copying TinyVGG from CNN Explainer
27     """
28     def __init__(self, input_shape: int,
29                 hidden_units: int,
30                 output_shape: int) -> None:
31         super().__init__()
32         self.conv_block_1 = nn.Sequential(
33             nn.Conv2d(in_channels=input_shape,
34                     out_channels=hidden_units,
35                     kernel_size=3,
36                     stride=1,
37                     padding=0),
38             nn.ReLU(),
39             nn.Conv2d(in_channels=hidden_units,
40                     out_channels=hidden_units,
41                     kernel_size=3,
42                     stride=1,
43                     padding=0),
44             nn.ReLU(),
45             nn.MaxPool2d(kernel_size=2,
46                         stride=2)
47         )
48         self.conv_block_2 = nn.Sequential(
49             nn.Conv2d(in_channels=hidden_units,
50                     out_channels=hidden_units,
51                     kernel_size=3,
52                     stride=1,
53                     padding=0),
54             nn.ReLU(),
55             nn.Conv2d(in_channels=hidden_units,
```

```

54         out_channels=hidden_units,
55         kernel_size=3,
56         stride=1,
57         padding=0),
58         nn.ReLU(),
59         nn.MaxPool2d(kernel_size=2,
60                     stride=2)
61     )
62     self.classifier = nn.Sequential(
63         nn.Flatten(),
64         nn.Linear(in_features=hidden_units*13*13,
65                 out_features=output_shape)
66     )
67
68     def forward(self, x):
69         x = self.conv_block_1(x)
70         x = self.conv_block_2(x)
71         x = self.classifier(x)
72         return x
73
74 model_0 = TinyVGG(input_shape=3,
75                  hidden_units=10,
76                  output_shape=len(class_names)).to(device)
77 model_0

```

```

TinyVGG(
  (conv_block_1): Sequential(
    (0): Conv2d(3, 10, kernel_size=(3, 3), stride=(1, 1))
    (1): ReLU()
    (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1))
    (3): ReLU()
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (conv_block_2): Sequential(
    (0): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1))
    (1): ReLU()
    (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1))
    (3): ReLU()
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (classifier): Sequential(
    (0): Flatten(start_dim=1, end_dim=-1)
    (1): Linear(in_features=1690, out_features=3, bias=True)
  )
)

```

2.2.6 Create train step, test step and train functions

```

1 # Create train_step()
2 def train_step(model: torch.nn.Module,
3               dataloader: torch.utils.data.DataLoader,

```

```

4         loss_fn: torch.nn.Module,
5         optimizer: torch.optim.Optimizer,
6         device=device):
7     # Put the model in train mode
8     model.train()
9
10    # Setup train loss and train accuracy values
11    train_loss, train_acc = 0, 0
12
13    # Loop through data loader data batches
14    for batch, (X, y) in enumerate(dataloader):
15        # Send data to the target device
16        X, y = X.to(device), y.to(device)
17
18        # 1. Forward pass
19        y_pred = model(X) # output model logits
20
21        # 2. Calculate the loss
22        loss = loss_fn(y_pred, y)
23        train_loss += loss.item()
24
25        # 3. Optimizer zero grad
26        optimizer.zero_grad()
27
28        # 4. Loss backward
29        loss.backward()
30
31        # 5. Optimizer step
32        optimizer.step()
33
34        # Calculate accuracy metric
35        y_pred_class = torch.argmax(torch.softmax(y_pred, dim=1), dim
36        =1)
37        train_acc += (y_pred_class==y).sum().item()/len(y_pred)
38
39    # Adjust metrics to get average loss and accuracy per batch
40    train_loss = train_loss / len(dataloader)
41    train_acc = train_acc / len(dataloader)
42    return train_loss, train_acc
43
44    # Create a test step
45    def test_step(model: torch.nn.Module,
46                  dataloader: torch.utils.data.DataLoader,
47                  loss_fn: torch.nn.Module,
48                  device=device):
49        # Put model in eval mode
50        model.eval()
51
52        # Setup test loss and test accuracy values
53        test_loss, test_acc = 0, 0
54
55        # Turn on inference mode
56        with torch.inference_mode():
57            # Loop through DataLoader batches
58            for batch, (X, y) in enumerate(dataloader):
59                # Send data to the target device
60                X, y = X.to(device), y.to(device)

```

```

60
61     # 1. Forward pass
62     test_pred_logits = model(X)
63
64     # 2. Calculate the loss
65     loss = loss_fn(test_pred_logits, y)
66     test_loss += loss.item()
67
68     # Calculate the accuracy
69     test_pred_labels = test_pred_logits.argmax(dim=1)
70     test_acc += ((test_pred_labels == y).sum().item()/len(
test_pred_labels))
71
72     # Adjust metrics to get average loss and accuracy per batch
73     test_loss = test_loss / len(dataloader)
74     test_acc = test_acc / len(dataloader)
75     return test_loss, test_acc
76
77 from tqdm.auto import tqdm
78
79 # 1. Create a train function that takes in various model parameters
+ optimizer + dataloaders + loss function
80 def train(model: torch.nn.Module,
81         train_dataloader: torch.utils.data.DataLoader,
82         test_dataloader: torch.utils.data.DataLoader,
83         optimizer: torch.optim.Optimizer,
84         loss_fn: torch.nn.Module,
85         epochs: int = 5,
86         device = device):
87
88     # 2. Create empty results dictionary
89     results = {"train_loss": [],
90             "train_acc": [],
91             "test_loss": [],
92             "test_acc": []}
93
94     # 3. Loop through training and testing steps for a number of
epochs
95     for epoch in range(epochs):
96         train_loss, train_acc = train_step(model=model,
97                                           dataloader=train_dataloader,
98                                           loss_fn=loss_fn,
99                                           optimizer=optimizer,
100                                           device=device)
101         test_loss, test_acc = test_step(model=model,
102                                       dataloader=test_dataloader,
103                                       loss_fn=loss_fn,
104                                       device=device)
105
106     # 5. Print out what's happening
107     print(f"Epoch: {epoch} | Train loss: {train_loss:.4f} | Train
acc: {train_acc:.4f} | Test loss: {test_loss:.4f} | Test acc: {
test_acc:.4f}")
108     results["train_loss"].append(train_loss)
109     results["train_acc"].append(train_acc)
110     results["test_loss"].append(test_loss)
111     results["test_acc"].append(test_acc)

```

```

112 # 6. Return the filled results at the end of the loop
113 return results

```

2.2.7 Train and evaluate model₀

```

1 # Set number of epochs
2 NUM_EPOCHS = 12
3
4 # Setup loss function and optimizer
5 loss_fn = nn.CrossEntropyLoss()
6 optimizer = torch.optim.Adam(params=model_0.parameters(),
7                               lr=0.001)
8
9 # Train model_0
10 model_0_results = train(model=model_0,
11                          train_dataloader=train_dataloader_simple,
12                          test_dataloader=test_dataloader_simple,
13                          optimizer=optimizer,
14                          loss_fn=loss_fn,
15                          epochs=NUM_EPOCHS)

```

```

Epoch: 0 | Train loss: 1.1063 | Train acc: 0.3047 | Test loss: 1.098321 | Test acc: 0.3011
Epoch: 1 | Train loss: 1.0998 | Train acc: 0.3281 | Test loss: 1.069704 | Test acc: 0.5417
Epoch: 2 | Train loss: 1.0869 | Train acc: 0.4883 | Test loss: 1.080744 | Test acc: 0.4924
Epoch: 3 | Train loss: 1.0843 | Train acc: 0.4023 | Test loss: 1.060744 | Test acc: 0.5833
Epoch: 4 | Train loss: 1.0662 | Train acc: 0.4102 | Test loss: 1.065383 | Test acc: 0.5644
Epoch: 5 | Train loss: 1.0298 | Train acc: 0.4414 | Test loss: 1.013412 | Test acc: 0.5426
Epoch: 6 | Train loss: 0.9809 | Train acc: 0.4180 | Test loss: 0.931092 | Test acc: 0.6146
Epoch: 7 | Train loss: 0.9571 | Train acc: 0.5859 | Test loss: 1.008202 | Test acc: 0.4744
Epoch: 8 | Train loss: 0.9266 | Train acc: 0.5898 | Test loss: 1.067889 | Test acc: 0.3324
Epoch: 9 | Train loss: 1.0050 | Train acc: 0.4609 | Test loss: 1.061566 | Test acc: 0.4044
Epoch: 10 | Train loss: 0.8879 | Train acc: 0.5195 | Test loss: 0.968867 | Test acc: 0.4830
Epoch: 11 | Train loss: 0.9317 | Train acc: 0.4453 | Test loss: 0.927652 | Test acc: 0.5436

```

2.2.8 Making a prediction on a custom image

```

1 # Download custom image
2 import requests
3
4 # Setup custom image path
5 custom_image_path = data_path / "04-pizza-dad.jpeg"
6
7 # Download the image if it doesn't already exist
8 if not custom_image_path.is_file():
9     with open(custom_image_path, "wb") as f:
10         request = requests.get("https://github.com/mrdbourke/pytorch-deep-learning/blob/main/images/04-pizza-dad.jpeg?raw=true")
11         f.write(request.content)
12 else:
13     print(f"{custom_image_path} already exists, skipping download...")

```

```

14
15 # Read in custom image
16 custom_image_uint8 = torchvision.io.read_image(custom_image_path)
17 plt.imshow(custom_image_uint8.permute(1, 2, 0))

```



```

1 # Create transform pipeline to resize image
2 from torchvision import transforms
3 custom_image_transform = transforms.Compose([
4     transforms.Resize(size=(64, 64)),
5 ])
6
7 # Transform target image
8 custom_image_transformed = custom_image_transform(custom_image)
9
10 plt.imshow(custom_image_transformed.permute(1, 2, 0))

```




Figure 1: Enter Caption

```

1 model_0.eval()
2 with torch.inference_mode():
3     custom_image_pred = model_1(custom_image_transformed.unsqueeze(0)
4     .to(device))
5
6 # Convert logits -> pred probs
7 custom_image_pred_probs = torch.softmax(custom_image_pred, dim=1)
8 print(custom_image_pred_probs)
9
10 # Convert pred prob -> pred labels
11 custom_image_pred_labels = torch.argmax(custom_image_pred_probs,
12     dim=1)
13 print(custom_image_pred_labels)
14 print(class_names[custom_image_pred_labels])

```

```

tensor([[0.3693, 0.3537, 0.2770]], device='cuda:0')
tensor([0], device='cuda:0')
pizza

```

3 Bibliography

<https://en.wikipedia.org/wiki/RegionBasedConvolutionalNeuralNetworks>
<https://towardsdatascience.com/review-fcn-semantic-segmentation-eb8c9b50d2d1>
<https://web.pdx.edu/~jduh/courses/Archive/geog481w07/Students/LudwigImageConvolution.pdf>
https://www.youtube.com/watch?v=V_xro1bcAuA